

Robust Federated Primal-Dual Learning for Android Malware Classification via Adversarial Robustness

M Mashreghi, *Staff, IEEE,*

Abstract—Abstract. The escalating frequency of Android malware attacks necessitates innovative approaches for detection and mitigation. Traditional machine learning (ML) and deep learning (DL) techniques face challenges due to the widespread distribution of Android devices, posing privacy concerns and operational overhead. Federated learning (FL) emerges as a promising solution, addressing privacy preservation and scalability issues. In this study, we explore alternative strategies to enhance federated learning techniques, encompassing non-id and id federated learning, FedAveraging (FedAvg), FedProx, and FedADMM. Additionally, we investigate the robustness of these approaches against adversarial attacks, specifically Projected Gradient Descent and Fast Gradient Sign Method. The evaluation, conducted on benchmark datasets (Melgenome, Drebin, Kronodroid, Tuandromd), unveils the scalability and privacy preservation of our approach. Notably, our proposed method outperforms traditional FedAvg in terms of accuracy, F1 score, AUC score, and FPR score for Android malware classification. This research contributes to the advancement of federated learning techniques and provides insights into their resilience against adversarial robustness, addressing the multifaceted challenges presented by Android malware attacks in an ever-evolving technological landscape.

Index Terms—Android Malware Classification, Federated Learning, Android Security, Distributed Machine Learning, Artificial Neural Network, Federated ADMM, Federated Prox, Fast Gradient Sign Method, Projected Gradient Descent, non-identical, Data Heterogeneity.

I. INTRODUCTION

IN contemporary times, Android has firmly established itself as one of the most widely used operating systems [2]. However, this ubiquity has come at a cost, with the proliferation of Android malware presenting a significant and escalating threat. Reports indicate that Android malware constitutes over 46% of all mobile malwares, marking a staggering 400% increase since 2010 [13]. Malicious software targeting Android devices has become a critical concern, especially given the unique features of these devices that make installing and using applications more accessible than on traditional computers. This heightened vulnerability not only poses challenges for individual users but also raises serious concerns for organizations.

The integration of machine learning (ML) and deep learning (DL) algorithms has become crucial in the domain of Android malware classification. Traditional ML/DL-based techniques face scalability issues in the current landscape marked by the rapid growth of mobile devices. Challenges such as decentralized data—where user-generated data is geographically

distributed—and the presence of sensitive information in the data further complicate traditional approaches.

To address these challenges, federated learning (FL) has emerged as a novel approach in malware detection and classification, offering enhanced security and efficient utilization of the Internet of Things (IoT) network system [5]. This paper embarks on a comprehensive exploration of federated models for cybersecurity and machine learning, dissecting the discussion into two key components. The first part elucidates federated learning (FL) and its application in IoT cybersecurity, while the second part focuses on cybersecurity for federated learning. The survey places significant emphasis on security approaches and addresses performance issues related to FL.

Noteworthy contributions in this domain include the Fed-IIoT architecture proposed in [6], integrating FL with an Android malware detection algorithm. The architecture incorporates dynamic poisoning attacks based on generative adversarial networks (GAN) and federated GAN to enhance robust collaboration training models. These models aim to circumvent anomalies in aggregation through a GAN network defense algorithm and adapt Byzantine defense algorithms, such as Krum and Median, on the server side.

In the context of Android malware classification, various innovative approaches have been proposed, incorporating data heterogeneity federated learning, FedAveraging, FedProx, and FedADMM. Non-id federated learning considers the scenario where participating devices have non-identical datasets, acknowledging the diverse data distributions in decentralized systems. On the other hand, id federated learning assumes identical datasets across participating devices, simplifying the federated learning process.

FedAveraging, a fundamental federated learning algorithm, creates a global model by aggregating local models from participating devices. FedProx extends this by incorporating a regularization term, mitigating issues related to non-identical data distributions. FedADMM, another federated learning optimization algorithm, focuses on decentralized optimization to solve the federated learning problem while minimizing direct communication between devices.

Furthermore, the paper explores the significance of considering adversarial attacks like Projected Gradient Descent (PGD) and Fast Gradient Sign Method (FGSM) to fortify the robustness and security of Android malware classification. PGD involves iteratively applying small perturbations to input data to deceive the model, while FGSM is a fast and computationally efficient method for crafting adversarial examples.

These enhancements contribute to the continued effective-

ness and adaptability of federated learning techniques, addressing the complex challenges posed by the escalating prevalence of Android malware attacks.

II. FUNDAMENTAL CONCEPTS

In this section we have provided the brief introduction about android malware and federated learning, FedProx, FedADMM, Robustness.

A. Android Malware

In the dynamic landscape of cyber threats, Android malware has emerged as a potent adversary, specifically crafted to exploit vulnerabilities in mobile devices running the Android operating system. Its proliferation has witnessed a significant surge, commanding a share exceeding 46% among various types of mobile malwares. Shockingly, there has been a fourfold increase in Android-based malware since 2010, highlighting the urgency of addressing this escalating menace. The evolving functionalities of Android devices, particularly in app installation and usage, introduce new challenges for the Android operating system in its defense against malware. This shifting paradigm necessitates continuous research endeavors.

B. Federated Learning

Federated learning, also known as collaborative learning, is a machine learning paradigm that diverges from traditional centralized techniques by training algorithms through multiple independent sessions, each utilizing its own dataset. This decentralized approach avoids merging local datasets into a single session and does not assume identical distribution of local data samples. Federated learning addresses critical issues such as data privacy, security, access rights, and heterogeneous data concerns. Its applications span diverse industries, including defense, telecommunications, Internet of Things (IoT), and pharmaceuticals.

The methodology involves training local models on local datasets and exchanging parameters, such as weights and biases, between nodes to create a global model shared by all. The key distinction from distributed learning lies in federated learning's focus on training on heterogeneous datasets. Clients involved may be less reliable due to factors like reliance on Wi-Fi and battery-powered systems. The mathematical formulation involves optimizing an objective function across nodes, aiming for consensus on model parameters.

Federated learning offers various setups, including centralized, decentralized, and heterogeneous approaches. Centralized federated learning involves a central server orchestrating the process, while decentralized federated learning allows nodes to coordinate without a central server. Heterogeneous federated learning accommodates diverse clients, such as mobile phones and IoT devices, addressing varying computation and communication capabilities.

The iterative nature of federated learning relies on rounds of client-server interactions, ensuring good task performance through a process of transmitting global model states, local

training, aggregation, and updating. Recent developments address asynchronicity and dynamically varying models. Non-IID data challenges the assumption of identical distribution across nodes, and techniques like data normalization are employed to mitigate the impact on training accuracy. Federated learning presents open questions regarding its preference over pooled data learning and concerns about device trustworthiness and the influence of malicious actors on learned models.

C. Robust optimization

Robust optimization serves as a critical mathematical framework designed to tackle the inherent uncertainties and fluctuations encountered in real-world systems. Widely applied in engineering, finance, and other fields sensitive to environmental changes and parameter uncertainties, its primary goal is to devise solutions that exhibit robust performance across a spectrum of potential scenarios. Unlike traditional optimization approaches, which assume precise knowledge of parameters and constraints, robust optimization recognizes the variability in these inputs, addressing the need for resilience in decision-making.

In classical optimization models, deterministic inputs are employed under the assumption of fixed and known conditions. However, practical situations often introduce variations, leading to suboptimal outcomes. Robust optimization responds to this challenge by considering a diverse range of potential input variations. Its strategic approach involves identifying and modeling uncertainties, integrating these variations into the optimization process, and formulating solutions capable of adapting to the unpredictability inherent in real-world data.

The core strength of robust optimization lies in its ability to strike a delicate balance between achieving optimal performance under nominal conditions and demonstrating adaptability to unforeseen variations. This makes it an invaluable tool in decision-making processes where uncertainty plays a substantial role, ensuring that solutions remain effective and reliable across a broad array of potential scenarios.

III. METHODOLOGY

In this section, the proposed approach details, the benchmark dataset description, experimental setup, and finally the performance metric are provided.

A. Federated Methods

1) *Federated Averaging (FedAvg)* : Federated Averaging (FedAvg) is a decentralized machine learning approach that allows model training across multiple devices or servers without exchanging raw data. It's a popular method in the field of federated learning, where the goal is to train a global model by aggregating updates from multiple local models.

Here's a high-level overview of how FedAvg works:

Algorithm 1 FedAvg

Input: N clients, each with a local dataset D_n and a local model w_n

Output: A global model w

- 1: Initialize the global model w_0
- 2: **for** $t = 1$ to T **do**
- 3: Randomly select a subset of N clients
- 4: **for** each client n in the subset **do**
- 5: Get model w_t^g from the server
- 6: Updating the local weights
- 7: Send the new local weights to the server
- 8: **end for**
- 9: Aggregate the local updates: $w_{t+1}^g = \frac{1}{N} \sum_{n=1}^N w_{n,t}^n$
- 10: Share w_{t+1}^g to all clients
- 11: **end for**

This process of local training, model update, and aggregation is repeated iteratively, allowing the global model to improve over time without the need to share raw data centrally. Federated learning is particularly useful in scenarios where data privacy and security are concerns, as the raw data never leaves the local devices.

2) *FedProx*: FedProx builds upon FedAvg and introduces a proximal term to the optimization objective. This proximal term addresses the challenge of non-identically distributed data among devices. Here's a simplified overview of FedProx:

Similar to FedAvg, devices compute the gradient of their local loss with respect to the model parameters. In addition to the local gradient, FedProx introduces a proximal term that penalizes deviations of the local model from the global model. Devices send both the local gradient and the proximal term to the central server. The central server aggregates these contributions and updates the global model. The proximal term in FedProx encourages devices to have models that are not only accurate on their local data but also similar to the global model. This helps mitigate issues arising from non-identically distributed data.

The proximal term in the FedProx strategy refers to a regularization term in the local model that is used on each client.

$$\min_w h_k(w; w_t^g) = F_k(w) + \frac{\mu}{2} \|w - w_t^g\|^2 \quad (1)$$

Where $F_k(w)$ is the loss function of k -th client which want to minimize itself, w are the local parameters of k -th client to minimize, w_t^g are the global parameters, and finally $h_k(w; w_t^g)$ is the objective function for Client k to minimize its weight with the regularization term.

Note FedAvg is a particular case of FedProx with $\mu = 0$. So, we just need to implement the code for FedProx, which we will be used also for FedAvg by setting the parameter $\mu = 0$

Algorithm 2 FedProx

Input: N clients, μ , each with a local dataset D_n and a local model w_n

Output: A global model w

- 1: Initialize the global model w_0
- 2: **for** $t = 1$ to T **do**
- 3: Randomly select a subset of N clients
- 4: **for** each client n in the subset **do**
- 5: Get model w_t^g from the server
- 6: Updating the local weights with the proximal term $\frac{\mu}{2} \|w - w_t^g\|^2$
- 7: Send the new local weights to the server
- 8: **end for**
- 9: Aggregate the local updates: $w_{t+1}^g = \frac{1}{N} \sum_{n=1}^N w_t^n$
- 10: Share w_{t+1}^g to all clients
- 11: **end for**

In summary, while FedAvg focuses on averaging local model updates, FedProx extends this approach by incorporating a proximal term to promote similarity between local and global models, making it more robust in scenarios with diverse or non-identical data distributions across devices.

3) *FedADMM*: FedADMM is a federated learning algorithm designed to address the challenges of communication efficiency, heterogeneity in client resources, non-i.i.d. data, and data privacy. It focuses on solving non-convex composite optimization problems with non-smooth regularizers.

The algorithm utilizes the concept of partial participation, where only a subset of clients participate in each round of communication. This addresses issues such as communication problems with a large number of clients or limited bandwidth in mobile phones.

FedADMM is based on the alternating direction method of multipliers (ADMM) and shares similarities with the FedDR algorithm, which combines the non-convex Douglas-Rachford splitting (DRS) algorithm with a randomized block-coordinate strategy. By leveraging the dual formulation of DRS, FedADMM handles the flexibility needed for federated learning.

Here's a high-level overview of the FedADMM algorithm: now we have:

now when we don't have $g(x)$ we can use this:

The contributions of FedADMM are as follows:

It proposes a new algorithm, FedADMM, based on the dual formulation of the problem, allowing partial participation and solving federated composite optimization problems. When the regularizer function is zero, FedADMM reduces to the FedPD algorithm but requires only partial participation. It establishes equivalence between FedDR and FedADMM, showing a one-to-one mapping between their iterates. It provides convergence guarantees for FedADMM using the equivalence established. FedADMM inherits desirable properties from FedDR, such as handling statistical and system heterogeneity and allowing inexact evaluation of users' proximal operators. It also expands the applicability to more general applications and problems with constraints.

Overall, FedADMM is a federated learning algorithm that leverages the ADMM framework, allowing partial participa-

Algorithm 2 Federated ADMM Algorithm (FedADMM)

```

1: Initialize  $x^0, \eta > 0, K$ , and tolerances  $\epsilon_{i,0} (i \in [n])$ .
2: Initialize the server with  $\bar{x}^0 = x^0$ 
3: Initialize all clients with  $z_i^0 = 0$  and  $\hat{x}_i^0 = \bar{x}^0 = x^0$ .
4: for  $k = 0, \dots, K$  do
5:   Randomly sample  $S_k \subseteq [n]$  with size  $S$ .
6:   ▷ Client side
7:   for each client  $i \in S_k$  do
8:     receive  $\bar{x}^k$  from the server.
9:      $x_i^{k+1} \approx \arg \min_{x_i} \mathcal{L}_i(x_i, \bar{x}^k, z_i^k)$ 
10:     $z_i^{k+1} = z_i^k + \eta(x_i^{k+1} - \bar{x}^k)$  ◇Dual updates
11:     $\hat{x}_i^{k+1} = x_i^{k+1} + \frac{1}{\eta} z_i^{k+1}$ 
12:    send  $\Delta \hat{x}_i^k = \hat{x}_i^{k+1} - \hat{x}_i^k$  back to the server
13:   end for
14:   ▷ Server side
15:   aggregation  $\bar{x}^{k+1} = \bar{x}^k + \frac{1}{n} \sum_{i \in S_k} \Delta \hat{x}_i^k$ 
16:   update  $\bar{x}^{k+1} = \text{prox}_{g/\eta}(\bar{x}^{k+1})$ 
17: end for

```

Fig. 1. FedADMM

$$\mathcal{L}_i(x_i, \bar{x}^k, z_i) = f_i(x_i) + g(\bar{x}^k) + \left\langle z_i^k, x_i - \bar{x}^k \right\rangle + \frac{\eta}{2} \|x_i - \bar{x}^k\|^2$$

Fig. 2. loss function of fedADMM

tion and addressing the challenges associated with communication efficiency and heterogeneity in distributed optimization.

B. Robust Optimization

1) *Fast Gradient Sign Method*: FGSM, which stands for Fast Gradient Sign Method, is a technique used in adversarial machine learning to generate adversarial examples. Adversarial examples are intentionally crafted inputs that are designed to mislead machine learning models.

The FGSM attack works by exploiting the gradients of a trained model with respect to its input to generate perturbations that are added to the original input. These perturbations are computed in the direction that maximizes the loss function of the model, causing misclassification or incorrect predictions.

Here's a step-by-step explanation of how FGSM works:

-1. Select a target input that you want to generate an adversarial example for. This input could be an image, a text sequence, or any other type of input that the model accepts.

-2. Pass the target input through the trained model and obtain the output probabilities or predictions. This step involves performing a forward pass through the model.

-3. Calculate the loss between the predicted output and the desired target class. The loss function quantifies the discrepancy between the predicted output and the target class, indicating how "wrong" the model's prediction is.

-4. Compute the gradients of the loss function with respect to the input. The gradients represent the sensitivity of the model's output to changes in the input. They provide information about how the input should be modified to maximize the loss function.

When $g \equiv 0$, the server-side steps 15-16 of FedADMM reduce to the single step:

$$\bar{x}^{k+1} = \bar{x}^{k+1} = \bar{x}^k + \frac{1}{n} \sum_{i \in S_k} \Delta \hat{x}_i^k = \frac{1}{n} \sum_{i=1}^n \hat{x}_i^{k+1}.$$

Fig. 3. FedADMM

-5. Determine the sign of the gradients and scale them based on a small constant value (usually denoted as epsilon, ϵ). The sign determines whether the perturbations will be added or subtracted from the original input, while the scaling factor controls the magnitude of the perturbations.

-6. Generate the adversarial example by adding the scaled perturbations to the original input. The perturbations are calculated as the sign of the gradients multiplied by the epsilon value.

-7. The resulting adversarial example is then fed back into the model, and the model's predictions are re-evaluated. Ideally, the modified input should cause the model to produce an incorrect or undesired output.

The FGSM attack is considered a "fast" method because it only requires a single backpropagation step to calculate the gradients and generate the adversarial example. However, it is worth noting that newer and more sophisticated attacks have been developed since FGSM was introduced, which can sometimes be more effective in generating adversarial examples.

2) *Projected Gradient Descent*: PGD, which stands for Projected Gradient Descent, is an iterative optimization method used in adversarial machine learning to generate stronger and more robust adversarial examples compared to FGSM. It is an extension of the FGSM attack that performs multiple iterations to refine the perturbations applied to the input.

Here's a step-by-step explanation of how PGD works:

Select a target input for which you want to generate an adversarial example.

Set an initial perturbation value, often denoted as "delta," which represents the maximum allowable perturbation for each pixel or feature in the input.

Initialize the adversarial example as a copy of the original input.

Perform a certain number of iterations or steps. In each iteration, the following steps are repeated:

a. Pass the adversarial example through the model and obtain the output probabilities or predictions.

b. Calculate the loss between the predicted output and the desired target class.

c. Compute the gradients of the loss function with respect to the input.

d. Determine the sign of the gradients and scale them based on the delta value.

e. Add the scaled perturbations to the current adversarial example.

f. Project the adversarial example back onto an ϵ -ball around the original input. This step ensures that the perturbations do not exceed a certain magnitude, preventing the adversarial example from being too perceptually different from the original

input. The projection restricts the perturbations to stay within a predefined range.

After the specified number of iterations, the final adversarial example is obtained.

The key difference between PGD and FGSM is the iterative nature of PGD. By performing multiple iterations, PGD explores the space of potential perturbations more thoroughly, allowing for better optimization and stronger adversarial examples. The projection step in PGD helps control the magnitude of the perturbations and ensures that the resulting adversarial example is within a reasonable range.

PGD is a commonly used attack method in adversarial machine learning due to its effectiveness in generating robust adversarial examples. By iteratively refining the perturbations, PGD can overcome certain defense mechanisms and increase the chances of misclassification or incorrect predictions by the targeted model.

At the end let take short summary to the formula:

The adversarial example X_{adv} in FGSM is generated as follows:

$$X_{adv} = X + \epsilon \cdot \text{sign}(\nabla_X J(X, y))$$

The PGD attack generates the adversarial example $X_{adv}^{(t+1)}$ iteratively using the following update rule:

$$X_{adv}^{(t+1)} = \text{Clip}_{X, \epsilon} \left(X_{adv}^{(t)} + \alpha \cdot \text{sign}(\nabla_X J(X_{adv}^{(t)}, y)) \right)$$

where:

- T is the number of iterations.
- $X_{adv}^{(t)}$ is the perturbed example at iteration t .
- α is the step size.
- $\text{Clip}_{X, \epsilon}(\cdot)$ ensures that the perturbed example stays within an epsilon ball around the original example X .
- $\nabla_X J(X_{adv}^{(t)}, y)$ is the gradient of the loss with respect to the perturbed example at iteration t .

C. Dataset Description

We have used four datasets for our approach which are publicly available. The brief description of those considered datasets are as follows:

1. **Malgenome:** This dataset contains features from 3799 app samples where 2539 are benign and 1260 are android malwares from Android malware genome project [10]. It contains a total of 215 features.

2. **Drebin:** This dataset contains features from 15036 app samples where 9476 are benign and 5560 are android malwares from Drebin project [12]. It also contains 215 features.

3. **Tunadromd:** This dataset [3] contains features from 4465 app samples where 903 are benign and 3565 are android malwares. It contains a total of 241 features.

4. **Kronodrid:** This dataset contains features from 78137 app samples where 36935 are benign and 41382 are android malwares [6]. It contains a total of 463 features.

TABLE I
DATASET INFORMATION

Dataset Name	No. of Samples	No. of Features	Class Labels
Malgenome	3799	215	Benign: 2539, Malware: 1260
Drebin	15036	215	Benign: 9476, Malware: 5560
Tunadromd	4465	241	Benign: 903, Malware: 3565
Kronodrid	78137	463	Benign: 36935, Malware: 41382

D. Experimental setup

– Machine configuration: Windows 64 bit OS, processor core-i7-10750H with RAM 24 Gb–2400MHz and 6Gb-Nvidia GTX-1660 graphics.

– Software development: The code is implemented in Python 3.11 with the help of Pytorch in the backend.

– Base classifier and optimizer: For the base classifier we have used 4-layer feed-forward network with 1st hidden layer contains 200 neurons 2nd hidden layer contains 100 and the 3rd hidden layer contains 50 neurons. The selection of layers is done by trail-error approach because there is approach to set it automatically. In hidden layers ReLU is the activation function and in the output layer sigmoid is the activation function. In the output layer sigmoid being the activation function is because we have considered binary classification (Benign or Malware). For training the neural network we have used Stochastic Gradient Descent (SGD) optimizer because its popularly used in federated approach.

– Parameter setup: The batch size of the train loader is 32, the epoche is 32. We have used 80-20 Hold-Out cross-validation technique to train and test the model and the overall performance comparison. The learning rate of SGD is 0.01.

– Source-code: The source code and implementation details of our proposed approach can be found in Github(needs to complete)

E. Performance Metric

The accuracy, F1-score, area under the ROC curve (AUC), and False Positive Rate (FPR) are metrics for evaluating classification models. For binary classification, the mathematical formulas are calculated in terms of positives and negatives:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{F1-score} = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

$$\text{Precision} = \frac{TP}{TP + \frac{1}{2} \cdot (FP + FN)}$$

$$\text{AUC} = \frac{\sum \text{Rank}(+) - | + | \cdot (| + | + 1) / 2}{| + | + | - |}$$

$$\text{FPR} = \frac{FP}{FP + TN}$$

Where:

TP = True Positives

TN = True Negatives

FP = False Positives

FN = False Negatives

$\text{Rank}(+) = \text{Sum of the ranks of all positively classified samples}$

$|+| = \text{Total number of positive samples}$

$|-| = \text{Total number of negative samples}$

IV. RESULT




Fig. 4. IID Data

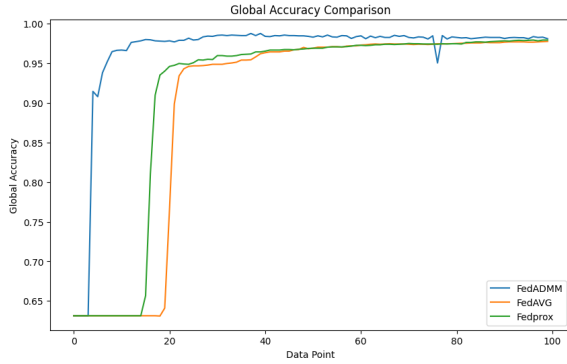


Fig. 5. IID PGD




Fig. 6. IID FGSM




Fig. 7. non IID Data




Fig. 8. non iid pgd

overall, we can see that FedADMM because of its loss function, Converge sooner than others. Also FedProx which has regularization term Converge sooner than FedAvg, in Non IID data we try to split data equally among the clients but we use Dirichlet Distribution for spread data from different classes to a client, as we can see in the result, it's like IID data, but the accuracy is a little less than IID data.

also if take attention to the PGD and FGSM attack we can see that PGD converge sooner than FGSM.




Fig. 9. non iid FGSM